Ranking-source comparison on synthetic development timetables

60 s per decision · 20 timetable groups · 3 paired solver seeds · no final-test or novelty claim

Each dot: median of 3 paired bill gains; all 3 valid positive-bill pairs required. Positive = lower bill. Panels use separate x scales.

v8 graph attention CatBoost (new, 17 Oct)

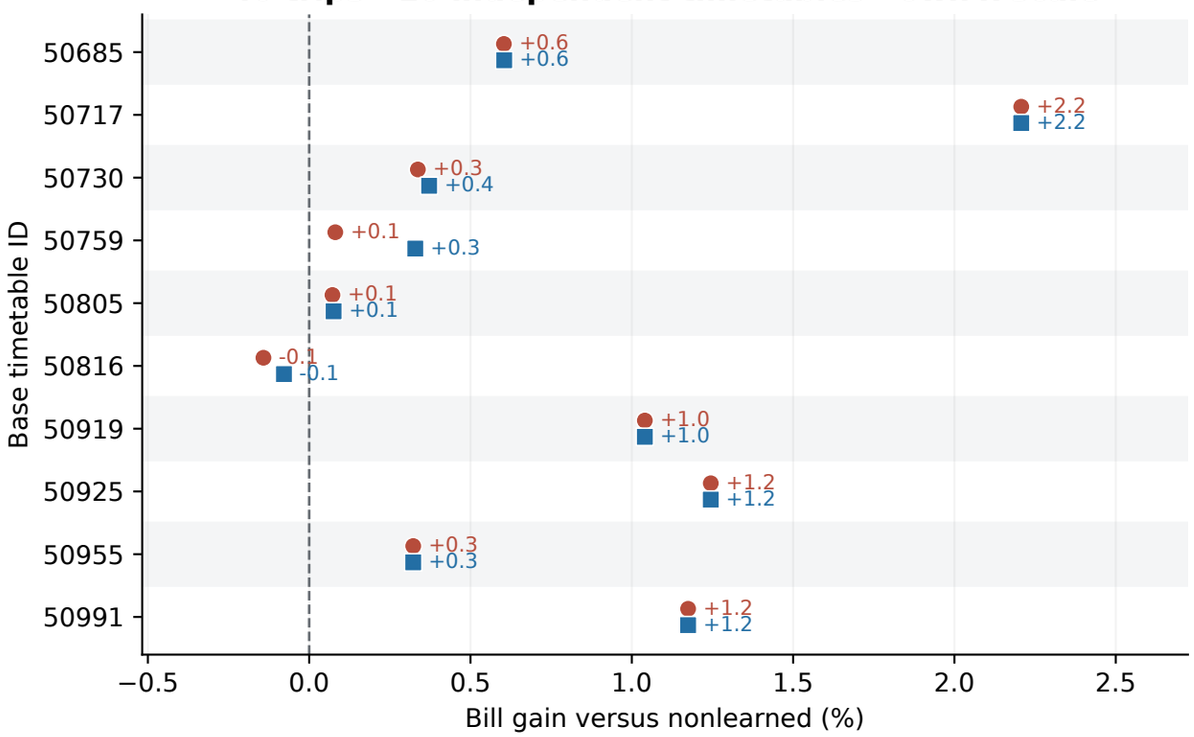

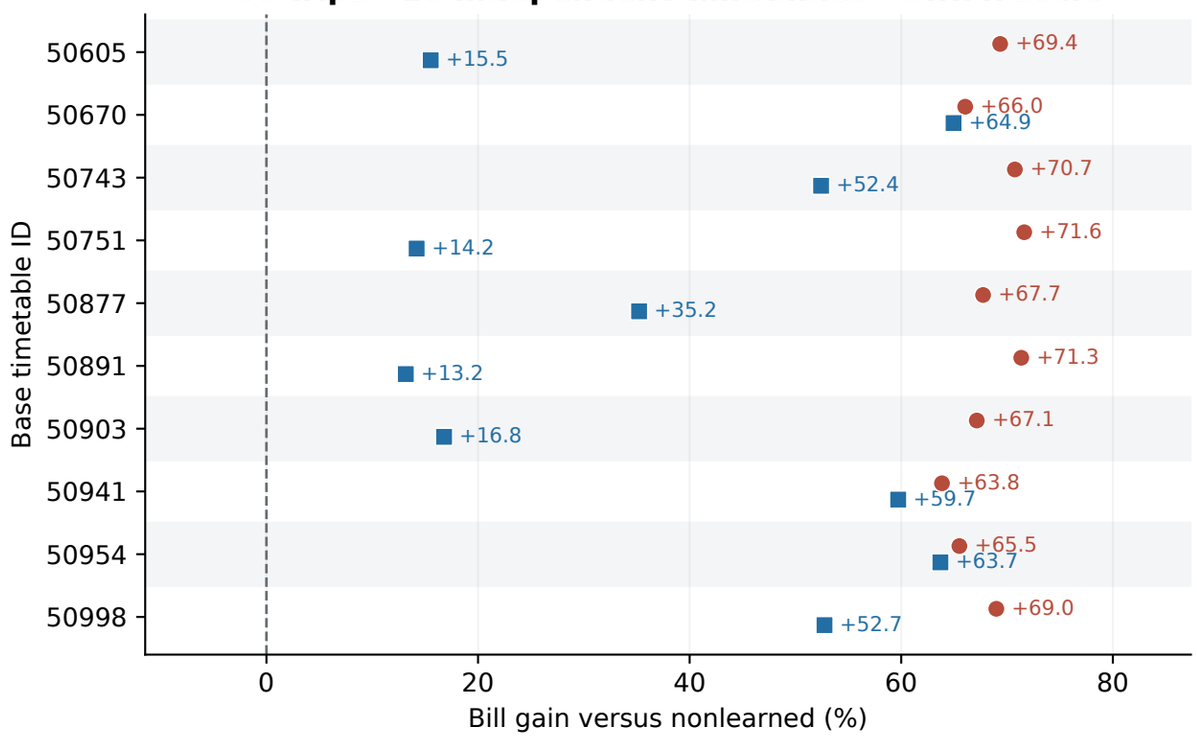

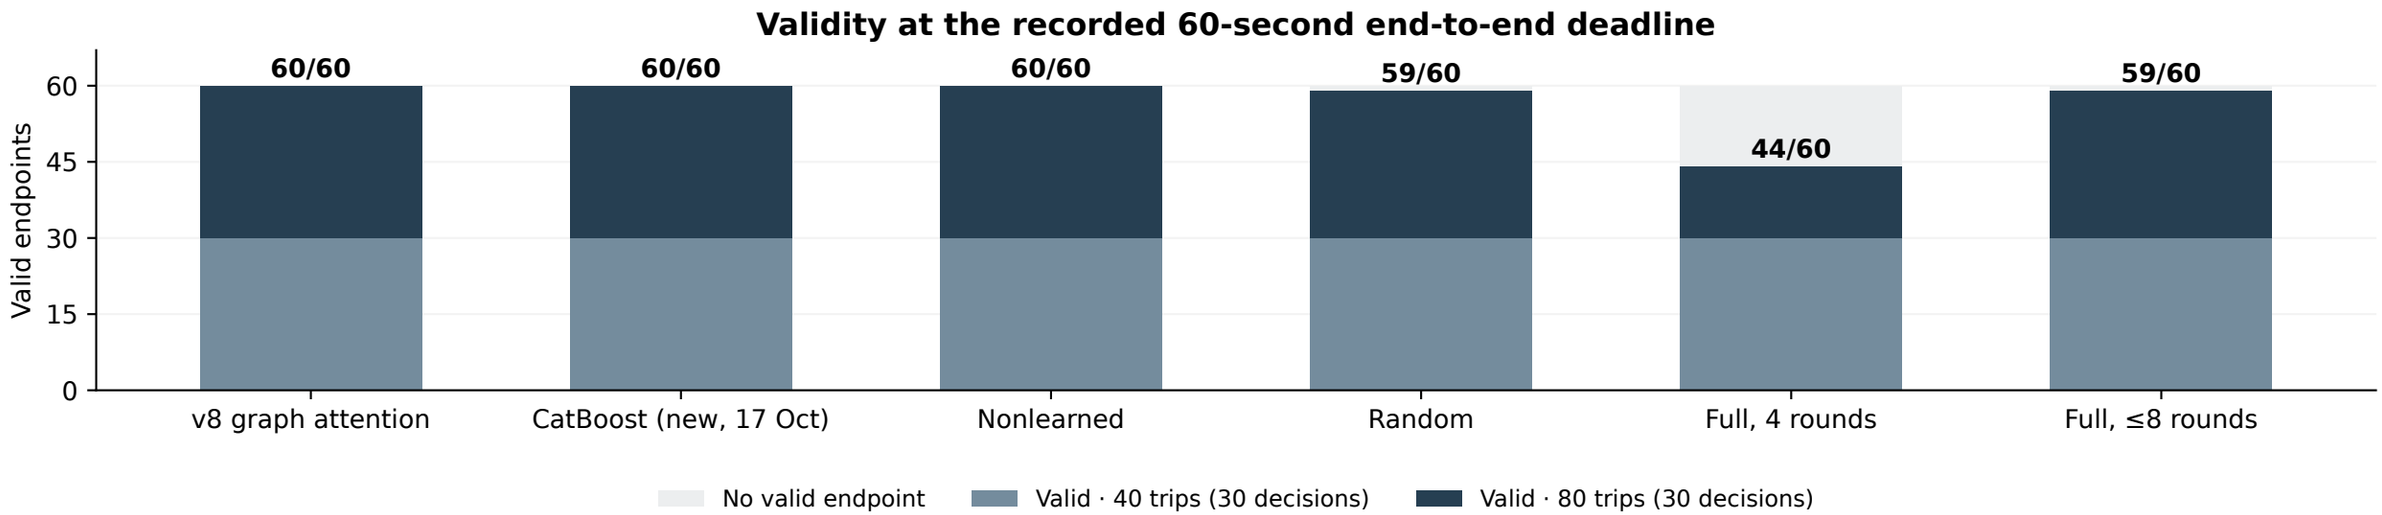

Stored supervisor validity and censoring are authoritative. Shaded bar remainder = no valid endpoint; seeds are repeats, not independent timetables.
CatBoost is newly trained (17 Oct 2026). Panels preserve size and timetable heterogeneity; no pooled headline median.